



Spark Program

CHAPTER 5:

SPARK STREAMING

Chapter Objectives

In this chapter, we will:

- Learn the special processing needs for the high-velocity data under Spark's streaming architecture
- Explore various streaming data sources

Introduction to Spark Streaming

- Spark Streaming is an extension of the core API
 - Provides high-throughput stream processing of live data
 - Built on Spark's fault-tolerant and highly scalable architecture
- Support for a wide variety of data sources
 - File systems, TCP Sockets, Kafka, Flume, Twitter, Kinesis, and ZeroMQ
 - Implement custom *Receivers* to integrate arbitrary data sources
- Streams can be processed using complex algorithms
 - Designed and implemented using:
 - Spark transformations and actions
 - Sliding window operations
 - The other extension APIs for SQL, Machine Learning, R, and GraphX
- Processed data can be pushed out to:
 - File systems, databases, live dashboards, etc.

Discretized Streams and RDDs

- Live input streams are divided into batches of data items
 - Known as *discretized streams*
 - Spark's high-level abstraction is called a DStream
- The batching interval is specified when the DStream is created
 - A new RDD is generated every interval for the batch of collected data
 - Spark operations are then applied to the RDD
 - Windowed DStream instances will contain multiple RDDs
 - A configurable number of historical RDDs from earlier batches
 - Spark operations are applied to their aggregate
- Input streams are obtained from Receivers
 - Built-in or custom classes for generating DStreams from data sources
 - Received data is reliably stored in Spark's memory for processing
- Spark provides two categories of supported streaming sources
 - Basic sources include file and socket streams
 - Advanced sources include Kafka, Flume, Kinesis, Twitter, etc.

Obtaining a DStream

- Applications must create a `StreamingContext`
 - Can use an existing `SparkContext` or `SparkConf` instance
 - At least two threads must be specified for the Worker/Executor and Receiver
 - Specify the batch interval

```
from pyspark import SparkContext
from pyspark.streaming import StreamingContext
sc = SparkContext("local[2]", "textStream")
sc.setLogLevel('Error')
ssc = StreamingContext(sc, 5)
lines = ssc.socketTextStream('localhost', 9999)
words = lines.flatMap(lambda line: line.split(' '))
pairs = words.map(lambda w: (w, 1))
wordCount = pairs.reduceByKey(lambda x, y : x + y)
wordCount.pprint()
ssc.start()
ssc.awaitTerminationOrTimeout(10000)
ssc.stop()
```

Setting the Batch Rate

- Applications should be able to process data as fast as it is being received
 - The batch rate is specified when the `StreamingContext` is instantiated
- The batch processing time should be less than the specified batch interval
 - Start with a longer than expected batch interval
 - Monitor the processing time and then reduce the batch interval
- Processing times can be monitored
 - Many useful streaming statistics are reported in the Web UI
 - The driver log files contain a “Total delay” entry
- Momentary delays may be acceptable
 - As long as the delay reduces back to an appropriate value
 - Caused by temporary increases in the data rate
- Sustained delays will accumulate over time
 - The application will not be able to keep up and may become unstable
- The level of processing parallelism can also be increased

DStream Transformations

- The following transformations are available on DStream instances: `map()`, `flatMap()`, `filter()`, `reduce()`, `reduceByKey()`, and `count()`
- `countByKey()`, `cogroup()`, and `join()`
 - These operate in exactly the same way as the standard RDD equivalents
- `transform()`
 - Applies any RDD to RDD operation not exposed by the DStream API
- `union()`
 - Combines multiple DStream instances of the same type

DStream Output Operations

- Output operations trigger the execution of any pending transformations
 - They behave like RDD actions
- Currently, DStream instances provide three save operations
 - Filenames are generated using the time with a prefix and optional suffix
 - Provided as arguments
 - `prefix-TIME_IN_MS.suffix`
- `saveAsTextFiles()`
- `saveAsObjectFiles()`
 - Save the DStream as serialized Java SequenceFile
 - Values only, keys are serialized as NullWritable
- `saveAsHadoopFiles()`
 - Requires a Hadoop OutputFormat
 - The DStream must contain (key, value) pairs
- `foreachRDD(func)`
 - Generic output operator that calls the `func` for each DStream

File Streams

- `fileStream()` monitors a file system directory
 - Called on a `StreamingContext` instance
 - Delegates to an underlying Hadoop `InputFormat`
- Files must be created in, or copied/moved to, the monitored directory
 - If required, existing files can be processed at start-up
 - File names starting with "." are ignored
 - A generic filename filter can be specified
- The `textFileStream()` method delegates to a `TextInputFormat`
- For full details, view the Spark Scala API documentation

```
ssc = StreamingContext(sc, 5)
lines = ssc.textFileStream('stream')
words = lines.flatMap(lambda line: line.split(' '))
pairs = words.map(lambda w: (w, 1))
wordCount = pairs.reduceByKey(lambda x, y : x + y).transform(lambda x :
x.sortBy(lambda x : -x[1]))
wordCount.pprint()
ssc.start()
ssc.awaitTerminationOrTimeout(10000)
ssc.stop()
```

DataFrames

- ➔ DStreams can be turned into DataFrames and then temporary views to make it easier to process them
- ➔ You need do a little trick to make the `SparkContext` object the same as the `StreamingContext`
 - The following lazy evaluated singleton instance will help do that

```
def getSparkSessionInstance(sparkConf):  
    if ("sparkSessionSingletonInstance" not in globals()):  
        globals()["sparkSessionSingletonInstance"] = \  
            SparkSession.builder \  
                .config(conf=sparkConf) \  
                .getOrCreate()  
    return globals()["sparkSessionSingletonInstance"]
```

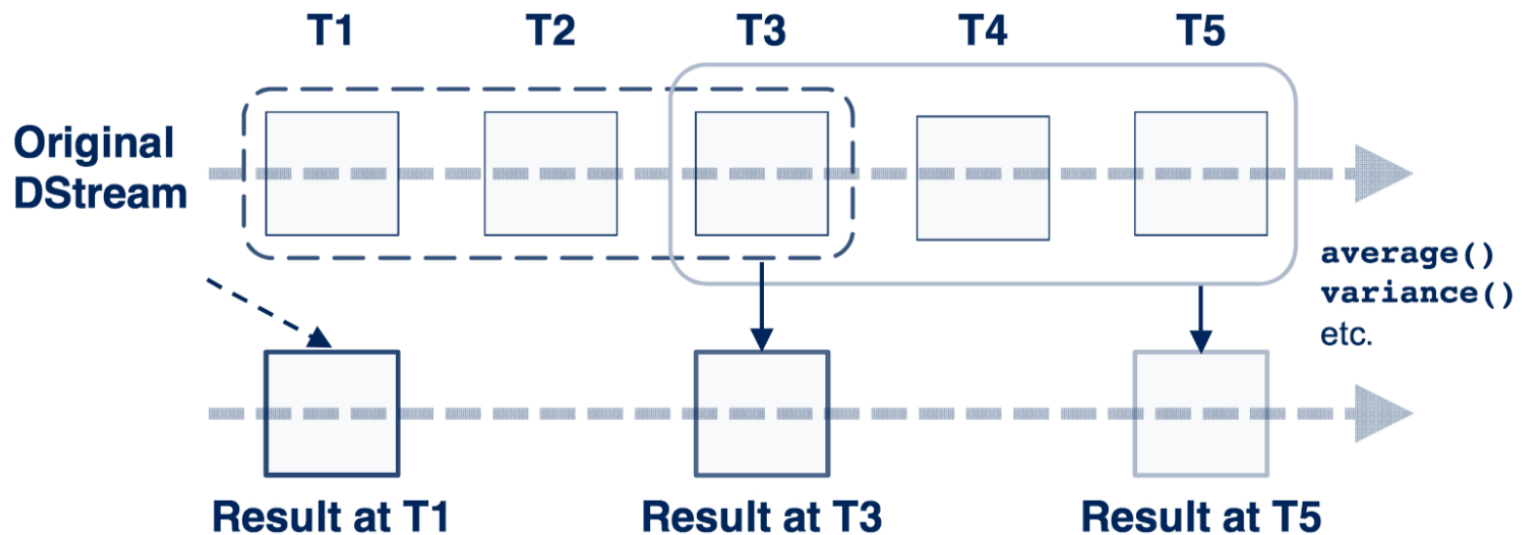
DataFrames (continued)

- ➔ To convert the DStream, you need to call the `foreachRDD` and pass it a function to call
- ➔ You need do a little trick to make the `SparkContext` object the same as the `StreamingContext`
 - The following lazy evaluated singleton instance will help do that

```
def process(time, rdd):  
    try:  
        spark = getSparkSessionInstance(rdd.context.getConf())  
        rdd1 = rdd.map(lambda x : x[1].split(',')) \  
                  .map(lambda x : (int(x[0]), float(x[1])))  
        df = spark.createDataFrame(rdd1, schema='id:int, amount:float')  
        df.createOrReplaceTempView('newdata')  
        join = spark.sql('select n.id, c.name, n.amount from newdata as n  
                          join codes as c on n.id = c.id')  
        join.show()  
    except:  
        print(rdd.collect())  
  
stream.foreachRDD(process)
```

Window Operations

- Spark streaming also provides an API for window computations
- As the window slides over a source DStream:
 - Operations are applied to the aggregate of RDDs that fall within the window
- Window operators must specify two parameters
 - Window length: the duration of the window
 - Sliding interval: the number of intervals to advance the window



Window Transformations

- Spark provides the following windowed transformations
 - Each requires a *window length* and *slide interval*
- `window()`
 - Returns DStream based on the *window length* and *slide interval*
- `countByWindow()`
 - Counts the number of elements in the window
- `countByValueAndWindow()`
 - Expects a DStream of (key, value) pairs
 - Returns a new DStream of (key, long) pairs in the window
- `reduceByWindow()` and `reduceByKeyAndWindow()`
 - Applies a reducing function to the values or (key, values) in the window
- The operation of `reduceByKeyAndWindow()` can be optimized
 - As a window slides, the reduced value can be calculated incrementally
 - An *inverse-reduce* function can be specified to remove old values
 - New values are then amalgamated by the reduce function

Advanced Streaming Sources

- Spark supports a number of advanced streaming sources
 - Not part of the core Spark API and require additional libraries
- Apache Kafka
 - A distributed publish-subscribe messaging system written in Scala
 - Designed to be fast, scalable, and robust
 - Originally developed by LinkedIn and became open source in 2011
- Apache Flume
 - Highly available distributed service for collecting and aggregating data
 - Designed to handle very large quantities of data
 - Originally developed by Cloudera
- Amazon Kinesis
 - Commercial EC2 service for collecting and processing stream-based data
 - Highly scalable and designed for used by real-time applications
 - Provide a Kinesis Client Library (KCL) under the Amazon Software License
- Twitter

Chapter Summary

In this chapter, we have:

- Learned the special processing needs for the high-velocity data under Spark's streaming architecture
- Explored various streaming data sources